

The Databricks Data Intelligence Platform

The Databricks Data Intelligence Platform is a multicloud unified, open analytics platform for building, deploying, sharing and maintaining enterprise-grade data, analytics and AI solutions at scale.



Compute

- Serverless** – Serverless compute is on-demand, automatically managed compute that scales based on your workload requirements.
- Classic** – Classic compute is provisioned compute resources that you create, configure and manage for your workloads.
- SQL Warehouse** – SQL warehouses are either classic or serverless and are optimized compute resources for specific use cases and advanced functionality.
- Lakebase** – A fully managed Postgres OLTP database engine that allows users to create and manage transactional databases in Databricks-managed storage, seamlessly integrate real-time OLTP workloads with your lakehouse, and support AI/BI workloads that require lower latency queries via reverse ETL patterns.
- Serverless GPU Compute** – On-demand, fully managed GPU infrastructure for running ML and AI workloads without manual resource management.


Data Sharing and Collaboration

- Delta Sharing** – The secure data sharing platform that lets users share data and AI assets within and outside Databricks.
- Clean Rooms** – Uses Delta Sharing and serverless compute to provide a secure and privacy-protecting environment where multiple parties can work together on sensitive enterprise data without direct access.
- Marketplace** – Open forum for exchanging third-party data products.
- Partner Connect** – Allows users to create trial accounts with select Databricks technology partners and connect their Databricks workspace to partner solutions from the Databricks UL.



Accelerators

- Lakebridge** – A unified tool that automates and simplifies the migration of legacy data warehouses and ETL workloads to Databricks SQL.
- Solution Accelerators** – Pre-built, customizable templates and best practices designed to help organizations quickly solve common industry and data analytics challenges on the Databricks Platform.



Data Engineering

- Lakeflow** – An end-to-end data engineering solution that empowers data engineers, software developers, SQL developers, analysts and data scientists to deliver high-quality data for downstream analytics, AI and operational applications.
- Connect** – Simplifies data ingestion with connectors popular for enterprise applications, databases, cloud storage, message buses and local files.
- Spark Declarative Pipelines** – A declarative framework that lowers complexity of building and managing efficient batch and streaming data pipelines by automatically orchestrating the execution of flows, sinks, streaming tables and materialized views.
- Jobs** – Provides reliable orchestration and production monitoring for any data and AI workload (notebooks, pipelines, managed connectors, SQL queries, machine learning training, model deployment and inference).
- Designer** – An AI-powered, visual no-code ETL builder that lets users create production-grade pipelines with drag-and-drop and natural language.
- Auto Loader** – Incrementally and efficiently processes new data files as they arrive in cloud storage.



Machine Learning

- MLflow 3.0** – Open source tool that simplifies and manages the complete machine learning lifecycle, including experiment tracking, model management and deployment, on Databricks and beyond.
- AutoML** – Automated machine learning capabilities that rapidly build, select and optimize models without extensive coding.
- Feature Store** – A centralized repository for storing, sharing and governing machine learning features to streamline model development and reuse.
- Machine Learning Runtime** – An optimized, up-to-date runtime environment tailored for scalable machine learning and AI workloads on Databricks.



Data Warehousing

- Databricks SQL (DBSQL)** – Collection of services that bring data warehousing capabilities and performance to users’ existing data lakes.
- SQL Editor** – Analyst-friendly user interface that blends a unified and streamlined ANSI-SQL authoring experience with multiple statement results, real-time collaboration and enhanced Databricks Assistant integrations.
- Lakehouse Federation** – A collection of features that enable users and systems to run queries against multiple data sources without needing to migrate all the data to a unified system.
- Query Federation** – With query federation, queries are pushed down to the foreign database using JDBC APIs and executed both in Databricks and using remote compute.
- Catalog Federation** – With catalog federation, users can directly access the foreign table in object storage using only Databricks compute.



Business Intelligence and Visualization

- Genie** – Allows business teams to interact with their data using natural language.
- AI/BI Dashboards** – Features AI-assisted authoring, an enhanced visualization library and a streamlined configuration experience for quickly transforming data into data visualizations and shareable reports.
- Databricks Apps** – Enables developers to build and deploy secure data and AI applications directly on the Databricks Platform, which eliminates the need for separate infrastructure.
- Databricks One** – A streamlined UI for workspace users focused on securely accessing published dashboards, Genie spaces and apps without creating new content
- Chat** – A unified, full-screen natural-language interface that uses Genie plus your existing dashboards and queries to answer data questions across your entire workspace.



Governance and Security

- Unity Catalog** – Unity Catalog is Databricks’ unified governance solution to centrally manage discovery, access controls, lineage, monitoring and sharing for all data and AI assets.
- Business Semantics** – Governed, reusable business metric definitions in Unity Catalog that ensure consistent, SQL-native access to trusted KPIs across platforms.
- Access Control** – Role-based and attribute-based policies that restrict or grant permissions to resources, ensuring secure and compliant data access.
- Data Lineage** – End-to-end tracking of data movement and transformations, providing transparency into how data is sourced, processed and consumed.
- Anomaly Detection** – Automated identification of unexpected patterns or unusual behavior in data or system performance.
- Governed Tags** – Standardized metadata labels applied to data assets to enforce compliance, track sensitivity and enable fine-grained governance.
- Data Classification** – Automated identification and labeling of sensitive data types to support compliance and secure data management.



Data Storage Layer

- Delta** – Delta is an open source table format that extends Parquet data files with a file-based transaction log for ACID transactions and scalable metadata handling.
- Apache Iceberg™** – Iceberg is an open source table format for analytics workloads, supporting features like schema evolution, time travel and hidden partitioning, along with an abstraction layer that enables ACID transactions on data stored in object storage.




AI

- Model Serving** – Scalable infrastructure for deploying and managing machine learning models as real-time REST APIs.
- Foundation Model Serving** – Managed deployment of large-scale foundation models (like LLMs) on Databricks for rapid generative AI integration.
- Vector Search** – A high-performance engine for storing and searching feature and embedding vectors to enable semantic retrieval and AI-powered search.
- Agent Framework** – A toolkit for building, orchestrating and managing AI agents that automate tasks across enterprise workflows.
- Agent Evaluation** – Automated tools and dashboards for benchmarking and improving the performance and reliability of AI agents.
- Gateway** – A centralized interface that manages secure, unified access to deployed AI and ML models and agents within an organization.
- AI Playground** – An interactive environment to experiment with models, prompts and AI components for rapid prototyping and learning.
- Agent Bricks** – Reusable, composable building blocks for quickly creating and deploying AI agent workflows in Databricks.
- Inference Tables** – Structured tables capturing inputs, outputs and metadata from model predictions to enable monitoring and analysis at scale.



Productivity

- Databricks Notebooks** – Primary tool for creating data science and machine learning workflows on Databricks; they provide real-time coauthoring in multiple languages, automatic versioning, and built-in data visualizations for developing code and presenting results.
- Genie Code** – An AI-powered pair programmer and support agent that helps you generate, optimize and troubleshoot code, queries and dashboards to accelerate your work.



Monitoring/Observability

- System Tables** – Built-in, queryable tables that provide operational metadata and insights into Databricks workspace activity.
- Lakehouse Monitoring** – Comprehensive tools and dashboards to observe, troubleshoot and optimize Lakehouse architecture health and workloads.
- Lakewatch** – An open, agentic SIEM that unifies security, IT and business data on the lakehouse to power AI-driven threat detection and response at petabyte scale and up to ~80% lower TCO than legacy SIEMs.

Cloud Providers

